

Sockets

Daniel Hagimont

IRIT/ENSEEIH
2 rue Charles Camichel - BP 7122
31071 TOULOUSE CEDEX 7

Daniel.Hagimont@enseeiht.fr
<http://hagimont.perso.enseeiht.fr>

1

The first part of this lecture is devoted to sockets.

What are sockets

- Interface for programming network communication
- Allow building client/server applications
 - Applications where a client program can make invocations to server programs with messages (requests) rather than shared data (memory or files)
 - Example: a web browser and a web server
- Not only client/server applications
 - Example: a streaming applications (VOD)

2

Sockets are a programming interface (API) for implementing message exchanges between processes which may run on different machines.

Such message exchanges are often used to implement distributed applications following the client-server model.

In this model, a server is a program running on one machine, which provides a service to some client programs running on other machines. A client may invoke this service by sending a message (called request) to the server program. Upon reception of a request, the server executes the treatments which correspond to the service, then it sends a message (called response) back to the client. The client is suspended after the emission of the request until reception of the response.

Notice that the request/response may include parameters/results.

A very popular example is the communication between a web browser (client) and a web server (server).

However, message exchanges can be used to implement other types of application, e.g. streaming applications like Video On Demand.

Two modes connected/not connected

- **Connected mode (TCP)**
 - Communication problems are handled automatically
 - Simple primitives for emission and reception
 - Costly connection management procedure
 - Stream of bytes: no message limits
- **Not connected mode (UDP)**
 - Light weight: less resource consumption
 - More efficient
 - Allow broadcast/multicast
 - All communication problems (packet loss) have to be handled by the application

3


Communication between processes can be performed following 2 modes :

- the connected mode corresponds to the use of the TCP communication protocol. We can create a connection between the client process and the server process. The connection is bi-directional (both the client and the server can send data on the connection). The communication mode is a stream of byte, i.e. there's no message limit. Communication problems (reemission of lost packets, blocking in case of buffer saturation) are automatically handled by TCP. The establishment of the connection is costly.

- the non connected mode corresponds to the use of the UDP communication protocol. There's no connection establishment anymore, nor handling of communication problems. It reduces resource consumption. A message of any size can be sent (it is split into IP packets, and reassembled on reception). There's no guarantee regarding message reception (it has to be handled by the application). Notice that UDP allows sending messages in multicast or broadcast (in general on a local network).

Sockets

- Network access interface
- Developed in Unix BSD
- @IP, #port, protocol (TCP, UDP, ...)



4

Sockets were initially developed in Unix BSD (Berkeley Software Distribution). They provide access to the network.

At the bottom layer (data link), machines (or rather network cards) are identified by a MAC address (e.g. an Ethernet address).

At the middle level (network), machines are identified by an IP address. ARP is the protocol which allows translating an IP address into a MAC address on a local network.

At the top level (transport), a process on one machine is identified by a couple @IP / #port, e.g. a web server is accessible on port 80 (default port) on a machine.

The socket API

- Socket creation: `socket(family, type, protocol)`
- Opening the dialog:
 - Client: `bind(..)`, `connect(...)`
 - Server: `bind(..)`, `listen(...)`, `accept(...)`
- Data transfer:
 - Connected mode: `read(...)`, `write(...)`, `send(...)`, `recv(...)`
 - Non-connected mode: `sendto(...)`, `recvfrom(...)`, `sendmsg(...)`, `recvmsg(...)`
- Closing the dialog:
 - `close(...)`, `shutdown(...)`

5

The socket API includes a set of functions in a programming language (initially C) for managing communication between processes.

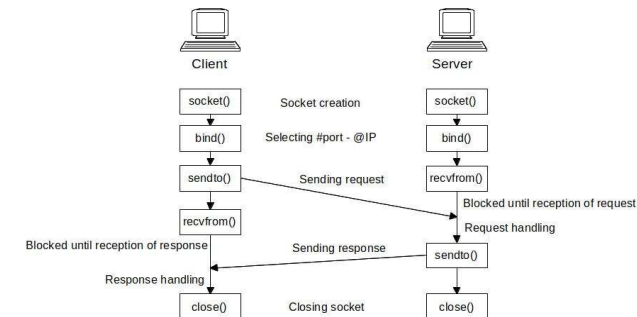
A socket is a file descriptor, similar to the file descriptors used to access files, except that writing or reading on such a descriptor sends or receives data to/from a remote process.

The socket API includes function for :

- creating a socket
- opening the dialog, i.e. initializing the socket (in connected or non-connected mode)
- transferring data (in connected or non-connected mode)
- closing the dialog

5

Client/Server in non-connected mode



6

We describe the schema of a request/response interaction between a client and a server. Here we consider its implementation with non connected sockets (UDP).

- both the client and the server create a socket with the `socket()` function which returns a file descriptor (fd, an index in the file descriptor table of the process). This fd is a parameter of all the following function calls.

- both the client and server call the `bind()` function which associates the socket with a local port of the machine (given as parameter). This port is the port used to receive messages (by the client or the server).

Generally, on the server side, this port is known in advance and given as parameter to `bind()`. The client knows this server port and communicates with the server identified with the IP address of the server and this server port. If the port is already used, `bind()` returns an error.

Generally, on the client side, the port given to `bind()` is 0, which means that `bind()` has to allocate a free port. This port is only used to receive responses.

- the client can call the `sendto()` function to send a message, giving as parameter the IP address and port of the target server process.

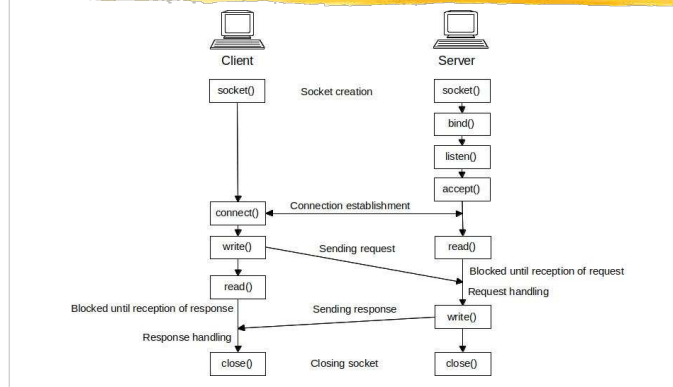
- the server can call the `recvfrom()` function to wait for a message. This function blocks until reception of a message. Upon reception, the message (request) is handled.

- the server can send a response with `sendto()`. The IP address and port of the client process (which sent the request) can be found in the request message.

- the client waits for the response using the `recvfrom()` function. Upon reception, the client can handle the response.

6

Client/Server in connected mode



Here we consider a request/response interaction with connected sockets (TCP).

- both the client and the server create a socket with the `socket()` function which returns a file descriptor (fd). This fd is a parameter of all the following function calls.

- on the server side

- `bind()` allows to associate the socket with a local port.

This port is generally known (e.g. port 80 for a web server)

- `listen()` allows to specify that the socket will be used to receive connection requests and how many connection requests can be pending

- `accept()` blocks until reception of a connection request from a client.

Upon reception of a connection request, `accept()` returns a **new socket** (a new fd) which is used by the server to send/receive on the established connection.

- on the client side

- `connect()` allows to send a connection request to the server, giving as parameter the IP address and port of the target server process. `connect()` includes a call to `bind()` (this is hidden).

After returning from `connect()`, the connection is established and the socket is used to send/receive.

What is important is the difference between the client and the server.

The client creates a socket, calls `connect()` and then use the socket to send/receive messages on that connection.

The server creates a socket, calls `bind()` and `accept()` and obtains a **NEW** socket for that connection with the client. The server may accept other connections with other clients and will obtain a different socket for each connection/client.

On a TCP connection, data may be sent/received with `write/read` functions on sockets (the same functions used to write/read data to/from a file).

7

socket() function

- `int socket(int family, int type, int protocol)`
- **family**
 - `AF_INET`: for Internet communications
 - `AF_UNIX`: for local communications
- **type or mode**
 - `SOCK_STREAM`: connected mode (TCP)
 - `SOCK_DGRAM`: non-connected mode (UDP)
 - `SOCK_RAW`: direct access to low layers (IP)
- **protocol :**
 - Protocol to use (different implementations can be installed)
 - 0 by default (standard)

8

We review the socket API in C.

`socket()` is the function which allows creating a socket.

`AF_UNIX` is used for local (to a machine) communications, while `AF_INET` is used for remote communications.


The type should be `SOCK_STREAM` for connected communication (TCP) and `SOCK_DGRAM` for non connected communication (UDP). Sockets can also be used in RAW mode (direct access to the IP level).

The protocol to be used should be 0 for default protocols (TCP, UDP), but could be different if other protocols are installed.

Notice that `socket()` returns an integer which is a file descriptor.

8

After call to socket()



9

This is a representation of the states of the client and server processes after a call to `socket()` on both sides.

On both sides, an entry in the file descriptor table was allocated for the socket.

bind() function

- `int bind(int sock_desc, struct sockaddr *my_@, int lg_@)`
- `sock_desc`: socket descriptor returned by `socket()`
- `my_@`: IP address and # port (local) that should be used
- Example (client or server):

```
int sd;
struct sockaddr_in my_address; // @IP, #port, mode

sd = socket(AF_INET, SOCK_STREAM, 0);
my_address.sin_family = AF_INET;
my_address.sin_port = 0; // let system choose a port
my_address.sin_addr.s_addr = INADDR_ANY;
                           // any network interface


bind(sd, (struct sockaddr *)&my_address, sizeof(my_address));
```

10

The `bind()` function is invoked on both sides. It creates the association between a socket and a local port.

- `sock_desc` is the fd of the socket
- `my_@` is a structure which describes initializations of the socket
 - `sin_port = 0` means that `bind()` should allocate a free port
 - `s_addr = INADDR_ANY` means `bind()` can use any network interface (in case there are several network interfaces (cards))
- `lg_@` is the size of the previous structure as it may differ depending on the OS

After call to bind()



We can already exchange messages in non-connected mode

11

This is a representation of the states of the client and server processes after a call to `bind()` on both sides.

On both sides, an socket in the file descriptor table is bound to a local port.

connect() function

- **int connect(int sock_desc, struct sockaddr * @_server, int lg_@)**
- **sock_desc**: socket descriptor returned by `socket()`
- **@_server**: IP address and # port of the remote server
- **Example of client:**

```
int sd;
struct sockaddr_in server; // @IP, #port, mode
struct hostent remote_host; // name et @IP


sd = socket(AF_INET, SOCK_STREAM, 0);
server.sin_family = AF_INET;
server.sin_port = htons(80);
remote_host = gethostbyname("www.enseiht.fr"); // DNS lookup
bcopy(remote_host->h_addr, (char *)&server.sin_addr,
      remote_host->hlength); // copy the address
connect(sd, (struct sockaddr *)&server, sizeof(server));
```

12

The `connect()` function is invoked on the client side. It sends a connection request to a remote server.

- `sock_desc` is the fd of the socket
- `@_server` is a structure which describes the remote server (@IP and port)
 - `sin_port` = the remote server port.
 - `htons` (host to network) is a function which converts the port number (13) from a host representation to a network representation. This comes from the fact that an integer may have different representations on different hardware (little indian, big indian)
 - `sin_addr` = the @IP of the remote server
 - `gethostbyname()` allows to obtain from DNS the IP from the machine name
 - The IP address is a structure which has to be copied into the `sin_addr` structure.
- `lg_@` is the size of the previous structure as it may differ depending on the OS

After call to connect()



13

This is a representation of the states of the client and server processes after a call to `connect()` on the client side.

A connection request has been sent from the client to the server.

listen() function

- **int listen(int sock_desc, int nbr)**
- **sock_desc:** socket descriptor returned by `socket()`
- **nbr:** maximum number of pending connections
- **Example of server:**

```
int sd;
struct sockaddr_in server; // @IP, #port, mode

sd = socket(AF_INET, SOCK_STREAM, 0);
server.sin_family = AF_INET;
server.sin_port = 0; // let system choose a port
server.sin_addr.s_addr = INADDR_ANY;
// any network interface
bind(sd, (struct sockaddr *)&server, sizeof(server));
listen(sd, 5);
```

14

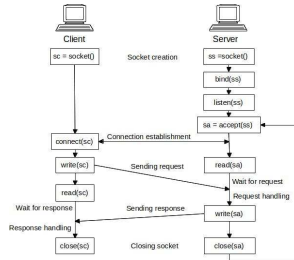
The `listen()` function is invoked on the server side to say that the socket will be used to receive connection requests and how many connection requests can be pending.

- `sock_desc` is the fd of the socket

- `nbr` is the number of tolerated pending connection requests (in a waiting queue). If the waiting queue is full, the connection from the client is rejected.

accept() function

- `int accept(int sock_desc, struct sockaddr *client, int lg_@)`
- **sock_desc**: socket descriptor receiving connection requests
- **client**: identity of the client which requested the connection
- **accept** returns the socket descriptor associated with the accepted connection




15

The `accept()` function is invoked on the server side. It blocks waiting for incoming connection requests. When a connection request is received, the blocked process is resumed and the function returns a new socket : the socket used to communicate with the client through the connection.

- `sock_desc` is the fd of the socket used to receive connection requests
- `client` is a structure which is updated with the identity (@IP, port) of the client who requested the connection.

After call to accept()



16

This is a representation of the states of the client and server processes after a connection has been accepted by the server.

In the server, a new socket (`#sock_connection`) was allocated and allows the server to communicate with the client through the connection.

Message emission/reception functions

- `int write(int sock_desc, char *buff, int lg_buff);`
- `int read(int sock_desc, char *buff, int lg_buff);`
- `int send(int sock_desc, char *buff, int lg_buff, int flag);`
- `int recv(int sock_desc, char *buff, int lg_buff, int flag);`
- `int sendto(int sock_desc, char *buff, int lg_buff, int flag, struct sockaddr *to, int lg_to);`
- `int recvfrom(int sock_desc, char *buff, int lg_buff, int flag, struct sockaddr *from, int lg_from);`
- `flag` : options to control transmission parameters (consult man)

17


Many functions are available for sending and receiving messages (the list here is not exhaustive).

The first four only take a socket and buffer as parameters, so they are used for the connection mode.

The last two take a `sockaddr` structure, allowing to specify the address (IP and port) we are sending to or to know the address of the sender we are receiving from. So they are used for the non connected mode.

Many functions have flags for controlling their behavior.

Communication



18

This figure illustrates communication on a TCP connection.

Both the client and the server can use read/write functions on the sockets associated with the connection. The connection is bi-directional.

A concurrent server

- After fork() the child inherits the father's descriptors

- Example of server:

```
int sd, nsd;
...
sd = socket(AF_INET, SOCK_STREAM, 0);
...
bind(sd, (struct sockaddr *)&server, sizeof(server));
listen(sd, 5);
while (!end) {
    nsd = accept(sd, ...);
    if (fork() == 0) {
        close(sd); // the child doesn't need the father's socket

        /* here we handle the connection with the client */

        close(nsd); // close the connection with the client
        exit(0); // death of the child
    }
    close(nsd); // the father doesn't need the socket of the connection
}
```

19

This is a typical example of concurrent server. The server is concurrent as a child process is created for each accepted connection.

The server creates a socket, binds it to a local port, and calls listen().

It then loops and waits for incoming connections (accept()). For each received connection, accept() returns a new socket (nsd). For this new connection, the server creates a process (fork()). The child process handles data received on this connection. The father process loops and waits for another connection.

Programming Socket in Java

- package **java.net**
 - **InetAddress**
 - **Socket**
 - **ServerSocket**
 - **DatagramSocket / DatagramPacket**

20

We now study the socket API in the Java environment.

Sockets in Java are provided by the java.net package.

The main classes are :

- InetAddress
- Socket and ServerSocket for TCP
- DatagramSocket and DatagramPacket for UDP

Using InetAddress (1)

```
import java.net.*;
public class Enseeiht1 {

    public static void main (String[] args) {
        try {
            InetAddress address =
                InetAddress.getByName("www.enseeiht.fr");
            System.out.println(address);
        } catch (UnknownHostException e) {
            System.out.println("cannot find www.enseeiht.fr");
        }
    }
}
```

21

InetAddress allows invoking the DNS, translating with `getByName()` a machine name into an IP address. It returns an `InetAddress` instance which includes the IP address.

Using InetAddress (2)

```
import java.net.*;
public class Enseeiht2 {

    public static void main (String[] args) {
        try {
            InetAddress a = InetAddress.getLocalHost();
            System.out.println(a.getHostName() + " / " +
                               a.getHostAddress());
        } catch (UnknownHostException e) {
            System.out.println("No access to my address");
        }
    }
}
```

22

InetAddress also allows to obtain the `InetAddress` of the local host. The returned `InetAddress` instance includes the machine name and its IP address.

Client socket and TCP connexion

```
try {
    Socket s = new Socket("www.enseeiht.fr",80);
    ...
} catch (UnknownHostException u) {
    System.out.println("Unknown host");
} catch (IOException e) {
    System.out.println("IO exception");
}
```

23

With TCP, a client can create a TCP connection with a target server (here www.enseeiht.fr), by creating an instance of the Socket class.

This operation corresponds to the calls in C of :

- socket()
- connect()

Reading/writing on a TCP connection

```
try {
    Socket s = new Socket ("www.enseeiht.fr",80);
    InputStream is = s.getInputStream();
    ...
    OutputStream os = s.getOutputStream();
    ...
} catch (Exception e) {
    System.err.println(e);
}
```

24

From this socket instance which is connected with the server, we can obtain 2 objects :

- an InputStream object which allows to read bytes
- an OutputStream object which allows to write bytes

With these objects, the client can send or receive data.

Server socket TCP connection

```
try {
    ServerSocket server = new ServerSocket(port);
    Socket s = server.accept();
    OutputStream os = s.getOutputStream();
    InputStream is = s.getInputStream();
    ...
} catch (IOException e) {
    System.err.println(e);
}
```

25

On the server side, the server can create a ServerSocket instance, giving a local port number as parameter. A ServerSocket instance is a socket for receiving connections. Therefore, this instantiation corresponds to the calls in C of :

- socket()
- bind()
- listen()

Then, the call of accept() on this instance blocks waiting for incoming connections. The process is resumed on connection reception, and accept() returns a Socket instance, which is the communication socket of the connection with the client. Like for the client side, we can obtain from this socket InputStream and OutputStream objects which provide communication methods.

Few words about classes for managing streams

- Suffix: type of stream
 - Stream of bytes (InputStream/OutputStream)
 - Stream of characters (Reader/Writer)
- Prefix: source or destination
 - ByteArray, File, Object ...
 - Buffered, LineNumber, ...
- <https://www.developer.com/java/data/understanding-byte-streams-and-character-streams-in-java.html>

26

InputStream and OutputStream are basic communication classes. They only allow to read and write bytes. They can be combined with many more elaborated classes.

The names of these classes are composed of a prefix and a suffix.

The suffix indicates the type of the stream

- suffix = InputStream or OutputStream for streams of bytes
- suffix = Reader or Writer for a streams of characters (unicode representation)

The prefix indicates the source or destination of the stream

- examples are File or Object

For instance :

- a FileInputStream allows reading bytes from a file
- a FileWriter allows writing characters to a file

Few words about classes for managing streams

	Streams for reading	Streams for writing
Character streams	BufferedReader CharArrayReader FileReader InputStreamReader LineNumberReader PipedReader PushbackReader StringReader	BufferedWriter CharArrayWriter FileWriter OutputStreamWriter PipedWriter StringWriter
Byte streams	BufferedInputStream ByteArrayInputStream DataInputStream FileInputStream ObjectInputStream PipedInputStream PushbackInputStream SequenceInputStream	BufferedOutputStream ByteArrayOutputStream DataOutputStream FileOutputStream ObjectOutputStream PipedOutputStream PrintStream

27

Here is a table of the different classes.

Few words about classes for managing streams

```
BufferedReader br = new BufferedReader(
    new InputStreamReader(socket.getInputStream()));
String s = br.readLine();
```

- InputStreamReader: converts a byte stream into a character stream
- BufferedReader: implements buffering

```
PrintWriter pred = new PrintWriter(
    new BufferedWriter(
        new OutputStreamWriter(
            socket.getOutputStream())));
```

- PrintWriter: formatted printing

28

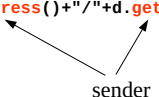
These classes can be combined (piped or chained) to obtain the desirable behavior.

In the first example, we obtain an InputStream from a socket. From this InputStream, we create a InputStreamReader which allows reading characters (Reader) from an InputStream (so it converts a stream of byte into a stream of characters). From this object, we create a BufferedReader, which provides buffering features like reading lines of characters.

In the second example, we obtain an OutputStream from a socket. From this OutputStream, we create a OutputStreamWriter which allows writing characters (Writer) to an OutputStream (so it converts a stream of characters into a stream of bytes). From this object, we create a BufferedWriter, which provides buffering features, and then a PrintWriter which provides formatted printing (like println()).

Reading on a UDP socket

```
try {
    int p = 9999;
    byte[] t = new byte[10];
    DatagramSocket s = new DatagramSocket(p);
    DatagramPacket d = new DatagramPacket(t, t.length);
    s.receive(d);
    String str = new String(d.getData(), 0, d.getLength());
    System.out.println(d.getAddress()+"-"+d.getPort()+"-"+str);
}
...
} catch (Exception e) {
    System.err.println(e);
}
```



29

A rapid look at programming UDP communication in Java.

On the receiving side, we create a DatagramSocket giving a local port number. It corresponds to the calls in C of : socket() and bind().

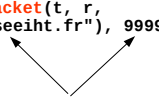
Then we can create a DatagramPacket giving an array of bytes.

Then

- receive() reads on the UDP socket and stores the data in the DatagramPacket
- getData() returns a byte array from the DatagramPacket (here we could have used the t variable)
- getLength() returns the size of the data actually received in the buffer
- getAddress() and getPort() return the address of the sender (IP and port), allowing to send a response.

Writing on a UDP socket

```
try {
    int p = 8888; // for receiving a response
    byte[] t = new byte[10];
    FileInputStream f = new FileInputStream("data.txt");
    int r = f.read(t);
    DatagramSocket s = new DatagramSocket(p);
    DatagramPacket d = new DatagramPacket(t, r,
        InetAddress.getByName("thor.enseeiht.fr"), 9999);
    s.send(d);
}
...
} catch (Exception e) {
    System.err.println(e);
}
```



30

On the sending side, we read in a byte array some data from a file.

We create a DatagramSocket giving a local port number.

Then we can create a DatagramPacket giving the array of bytes containing the data to send (r is the size of the data we read from the file), and also giving the destination (an InetAddress for the remote machine and a port from that machine). We send the packet with send().

A full example: TCP + serialization + threads

Passing an object (by value) with serialization

The object to be passed:

```
public class Person implements Serializable {
    String firstname;
    String lastname;
    int age ;
    public Person(String firstname, String lastname, int age) {
        this.firstname = firstname;
        this.lastname = lastname;
        this.age = age;
    }
    public String toString() {
        return this.firstname+" "+this.lastname+" "+this.age;
    }
}
```

31

Here is an example of client server communication with TCP, with the creation of a thread in the server on connection reception, and with an object passed with serialization.

Serialization is a Java mechanism which allows an instance to be copied between remote hosts (e.g. from a client to a server). The instance is translated into a byte array on the source machine and the instance is reconstructed on the destination machine. Serialization applies recursively, meaning that instances referenced (by a field) from one serialized instance are also serialized (so we can serialize a graph of objects). To enable serialization, a class must implement the Serializable interface. Notice that a serializable class should not include references to non serializable objects (e.g. a system resource like Thread or Socket).

Here we describe a serializable class (Person) that we will use to demonstrate the transfer (copy) of an instance on a TCP connection.

A full example: TCP + serialization + threads

The client

```
public class Client {
    public static void main (String[] str) {
        try {
            Socket csock = new Socket("localhost",9999);
            ObjectOutputStream oos = new ObjectOutputStream (
                csock.getOutputStream());
            oos.writeObject(new Person("Dan","Hagi",55));
            csock.close();
        } catch (Exception e) {
            System.out.println("An error has occurred ...");
        }
    }
}
```

32

Here is the client side of the TCP example.

The main() method :

- creates a Socket which connects to a server located at localhost/9999
- from the OutputStream of the socket, it creates an ObjectOutputStream, which allows writing objects to an OutputStream. This ObjectOutputStream (oos) serializes objects and sends the data on the connection.
- writes a Person instance on oos. The instance is then serialized.
- finally closes the socket

A full example: TCP + serialization + threads

The server

```
public class Server {
    public static void main (String[] str) {
        try {
            ServerSocket ss;
            int port = 9999;
            ss = new ServerSocket(port);
            System.out.println("Server ready ...");
            while (true) {
                Slave sl = new Slave(ss.accept());
                sl.start();
            }
        } catch (Exception e) {
            System.out.println("An error has occurred ...");
        }
    }
}
```

33

Here is the server side of the TCP example.

The main() method :

- creates a ServerSocket bound to local port 9999
- then it loops on connection reception
 - accept() blocks and when resumed by a connection reception, it returns a Socket instance.
 - it creates a Slave instance (giving it a reference to the Socket instance)
 - Slave is a class which implements a thread (explained next slide).
 - the thread is started

A full example: TCP + serialization + threads

The slave

```
public class Slave extends Thread {
    Socket ssock;
    public Slave(Socket s) {
        this.ssock = s;
    }
    public void run() {
        try {
            ObjectInputStream ois = new ObjectInputStream(
                ssock.getInputStream());
            Person v = (Person)ois.readObject();
            System.out.println("Received person: "+ v.toString());
            ssock.close();
        } catch (Exception e) {
            System.out.println("An error has occurred ...");
        }
    }
}
```

34

A way to program a thread is to implement a class which inherits from the Thread class. This class MUST implement the run() method which is invoked when the thread starts. NB : a thread is started with the start() method, not the run() method.

Here, the Slave class :

- inherits from Thread
- has a constructor to receive the socket it has to deal with
- implements a run() method which
 - creates an ObjectInputStream instance (ois) for reading objects from the stream of the socket. This ObjectInputStream instance reads data from the stream of the socket and deserializes the received objects.
 - reads an object on ois (the instance is deserialized) and casts it to Person (it is supposed to be a Person)

Conclusion

- Programming with sockets
 - Quite simple
 - Allow fine-grained control over exchanges messages
 - Basic, can be verbose and error prone
- Higher level paradigms
 - Remote procedure/method invocation
 - Message oriented middleware / persistent messages
 -

Many tutorials about socket programming on the Web ...

Example : https://www.tutorialspoint.com/java/java_networking.htm

35

In conclusion, programming with sockets is more or less simple (complex with C, simple in Java). It allows to do everything regarding distribution.

But for complex applications, even in Java, it may be really error prone.

This is why higher level programming paradigms were proposed.

In the next lectures, we will study some of them (remote invocation and message middleware).

Notice that many tutorials are available on the net for socket programming.

Client-server model

Daniel Hagimont

IRIT/ENSEEIH
2 rue Charles Camichel - BP 7122
31071 TOULOUSE CEDEX 7


Daniel.Hagimont@enseeiht.fr
<http://hagimont.perso.enseeiht.fr>

1

This lecture is about the client server model. It reviews the concepts and illustrates them with its instantiation in the Java environment, with Remote Method Invocation (RMI).

Client-server model based on message passing

- Two exchanged messages (at least)
 - The first message corresponds to the request. It includes the parameters of the request.
 - The second message corresponds to the response. It includes the result parameters from the response.



2

Client-server interactions can be implemented with message passing (using sockets).

You then have at least 2 messages exchanged for such an interaction.

The first message corresponds to the request, including parameters, and the second message corresponds to the response, including result parameters.

The client's execution is suspended after sending of the request, until reception of the response.

We can observe that such an interaction looks like a procedure call, except that the caller (client) and the callee (server) are located on different machines.

Remote Procedure Call (RPC) Principles

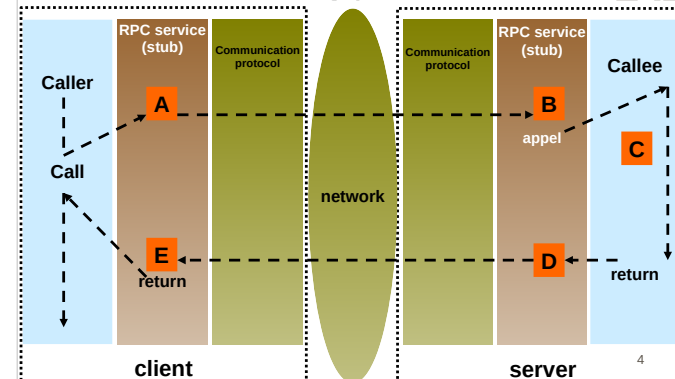
- Generating most of the code
 - Emission and reception of messages
 - Detection and re-emission of lost messages
- Objectives: the developer should be able to program the application without the burden to deal with messages

3

We call RPC (Remote Procedure Call) a tool which simplifies the development of applications relying on such client-server interactions, by generating the code which implements message exchanges (requests and responses). The idea is that all this code can be generated from a description of the interface of the procedure (which can be called on the server from a remote client).

The objective is to allow the developer to program and test his application as if it was centralized (executed on one machine) without the burden to deal with message exchanges. The code enabling the application to be distributed can be generated and the code of the application is kept simple.

RPC [Birrel & Nelson 84] Implementation principle



4

This is the general principle of RPC tools.

In blue, you have 2 code segments, the caller in the client which invokes (call) a service, and the callee in the server which provides the service.

In a centralized environment, the call would be a simple procedure call between the caller and the callee.

The principle of a RPC tool is to generate 2 code segments (in brown) called the client stub (left) and the server stub (right).

The client stub represents the service on the client machine and gives the illusion that the service is local (can be invoked locally with a simple procedure call). The client stub implements the same procedure as the server in order to give this illusion. A call to the procedure in the client stub creates and sends a request message to the server. The server stub receives and transforms request messages into local procedure calls on the server machine.

RPC (point A) Implementation principle

- On the caller side
 - The client makes a procedural call to the client stub
 - The parameters of the procedure are passed to the stub
 - At point A
 - The stub collects the parameters and assembles a message including the parameters (parameter marshalling)
 - An identifier is generated for the RPC call and included in the message
 - A watchdog timer is initialized
 - Problem: how to obtain the address of the server (a naming service registers procedures/servers)
 - The stub transmits the message to the transport protocol for emission on the network

5

On the caller side, the client performs a procedure call (invoking the service) as if the service was local to the client machine. Notice that the client stub implements the same procedure as the server, but the implementation of that procedure is different.

At point A, the client stub is called and receives the parameters from the procedure call. It assembles a request message which includes these parameters (this step is called parameter marshalling). An identifier for this RPC call is generated and included in the request message. This identifier allows to detect on the server side the reception of 2 requests for the same call (if the message is supposed to be lost and re-emitted).

A watchdog timer is initialized. It wakes up after a given time. If we don't receive a response before the wakeup, we consider that the request was lost and the request is re-emitted.

One problem here is to get the address (IP/port) of the server process for sending requests. Generally a naming service allows to register available procedures and their addresses.

The stub can then send the request message with the communication protocol (generally UDP as the data to be transmitted is not large).

The client is then suspended, waiting for the response message.

RPC (points B et C) Implementation principle

- On the callee side
 - The transport protocol delivers the message to the RPC service (server stub)
 - At point B
 - The server stub disassembles the parameters (parameter unmarshalling)
 - The RPC identifier is registered
 - The call is then transmitted to the remote procedure which is executed (point C)
 - The return from the procedure returns back to the server stub which receives the result parameters (point D)

6

On the callee side, the communication protocol delivers the request message to the server stub. At point B, the server stub disassembles the parameters of the call (this step is called parameter unmarshalling). The RPC identifier is registered to detect redundant requests for the same call.

The call is then reproduced, i.e. the procedure to be called in the callee is actually called (point C). This is a normal procedure call. On return, the procedure returns back (point D) to the server stub (with some result parameters).

RPC (point D) Implementation principle

- On the callee side
 - At point D
 - The result parameters are assembled in a message
 - Another watchdog timer is initialized
 - The server stub transmits the message to the transport protocol for emission on the network

7

At point D, the server stub assembles the result parameters in a response message.

Another watchdog timer is initialized. It wakes up after a given time. If we don't receive an acknowledgment from the client (that the response was received) before the wakeup, we consider that the response was lost and the response is re-emitted.

The server stub can then send the response with the communication protocol.

RPC (point E) Implementation principle

- On the caller side
 - The transport protocol delivers the response message to the RPC service (client stub)
 - At point E
 - The client stub disassembles the result parameters (parameter unmarshalling)
 - The watchdog timer created at point A is disabled
 - An acknowledgment message with the RPC identifier is sent to the server stub (the watchdog timer created at point D can be disabled)
 - The result parameters are transmitted to the caller with a procedure return

8

On the caller side, the communication protocol delivers the response message to the client stub.

At point E, the client stub disassembles the result parameters (parameter unmarshalling).

The watchdog timer created at point A can be disabled.

An acknowledgment message with the RPC identifier is sent to the server stub (the watchdog timer created at point D can be disabled).

The result parameters are transmitted to the caller with a procedure return.

RPC Role of stubs

Client stub

- It is the procedure which interfaces with the client
 - Receives the call locally
 - Transforms it into a remote call with a sent message
 - Receives results in a message
 - Returns results with a normal procedure return

Server stub

- It is the procedure on the server node
 - Receives the call as a message
 - Performs the procedure call on the server node
 - Receives the results of the call locally
 - Transmits the results remotely as a message

9

We summarize here the role of the client stub and the server stub.

RPC Message loss

- On the client side
 - If the watchdog expires
 - Re-emission of the message (with the same RPC identifier)
 - Abandon after N attempts
- On the server side
 - If the watchdog expires
 - Or if we receive a message with a known RPC identifier
 - Re-emission of the response message
 - Abandon after N attempts
- On the client side
 - If we receive a message with a known RPC identifier
 - Re-emission of the acknowledgment message

10

We provide here a global view of the handling of message loss.

On the client side, we created a watchdog before sending the request. If this watchdog expires, we can suppose that the request was lost and we re-send the request with the same RPC identifier (we re-initialize the watchdog before sending). We abandon after N attempts, assuming that the network is down.

On the server side, we create a watchdog before sending the response. As previously, we re-send the response if the watchdog expires. Another case on the server side is when we receive a request with a known RPC identifier (requests are logged). This means that we already received this request and the procedure was called and the response sent, but the response was lost. Then we re-send the response. As previously, we abandon after N attempts.

Finally, on the client side, if we receive a response with a known RPC identifier (response are logged), i.e. a response that we already received, it means that the acknowledgment sent to the server was lost and we re-send it.

RPC Problems

- Failure handling
 - Network or server congestion
 - The response arrives too late (critical systems)
 - The client crashes during the request handling on the server
 - The server crashes during the handling of the request
 - Failure of the communication system
 - What guarantees ?
- Security problems
 - Client authentication
 - Server authentication
 - Privacy of exchanges
- Performance
- Designation
- Practical aspects
 - Adaptation to heterogeneity conditions (protocols, languages, hardware)

11

Many other problems can be handled by RPC systems.

The handling of failures covers many types of failure :

- dealing with network or server congestion. Messages may be re-emitted, but redundant messages must be managed. In a real-time system, the execution time of a procedure is specified and the procedure should return an error if the deadline is not respected.

- dealing with the crash of the client or the server during the handling of the request, or the failure of the communication system. The system should provide guarantees (e.g. transactional behavior).

A RPC tool may also integrate security features, like authentication and encryption of exchanges.

Many other aspects were also considered :

- performance of RPC, especially the optimization when the client and server processes are on the same machine, or on the same LAN.

- designation : different designation scheme can be provided, for identifying the target (process) of call.

- heterogeneity : a lot of work was done to enable heterogeneity (of languages, OS ...) between the caller and callee (see CORBA).

RPC IDL : interface specification

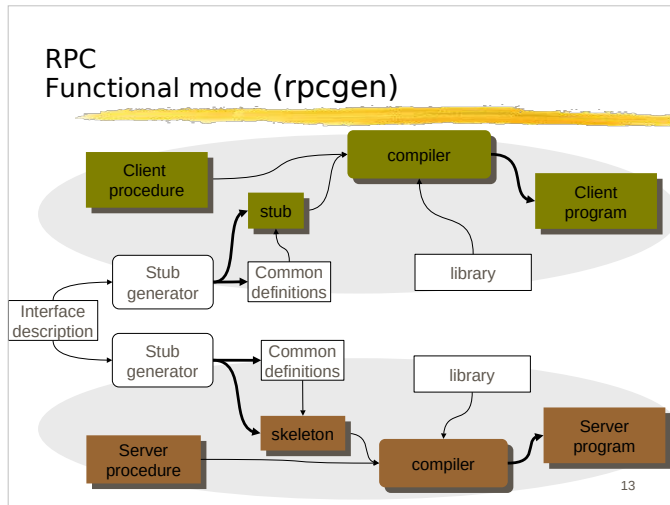
- Use of an interface description language (IDL)
 - Specification which is common to the client and the server
 - Definition of parameter types et natures (IN, OUT, IN-OUT)
- Use of the IDL description to generate:
 - The client stub (also called proxy or stub)
 - The server stub (also called skeleton)

12

Generally, a RPC tool generates stubs from the specification of the interface of the procedure which can be called remotely.

An IDL (Interface Description Language) is a simple language for describing the interface of a procedure which can be called through a RPC system. It simply allows describing the signature of the procedure, including the type of the parameters (data structures).

Such a specification allows to generate the client stub (sometimes called proxy or simply stub) and the server stub (often called skeleton).



rpcgen is one of the first RPC tools which was available in a Unix/C environment.

From the interface description (expressed with the IDL), a stub generator generates both the sub and skeleton.

On the client side, the client procedure (caller) is compiled with the stub in order to obtain an executable binary (client program).

On the server side, the server procedure (callee) is compiled with the skeleton in order to obtain a executable binary (server program).

These 2 binaries can be installed on different machines and executed.

Java Remote Method Invocation RMI

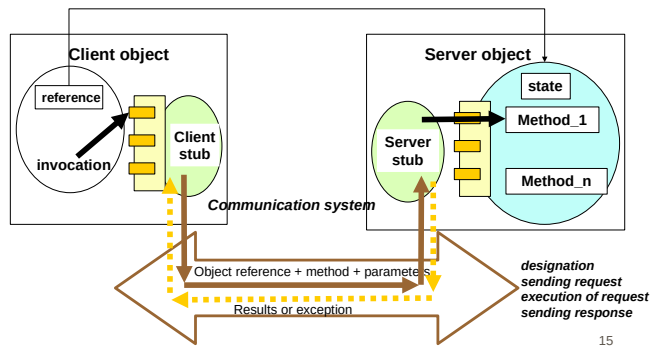
- An object based RPC integrated within Java
- Interaction between objects located in different address spaces (*Java Virtual Machines* - JVM) on remote machines
- Easy to use: a remote object is invoked as if it was local

14

Java RMI (Remote Method Invocation) is an example of implementation of a RPC tool integrated in a language environment (here Java).

It allows the invocation of methods on instances located on remote machines (in a remote JVM). Such a remote method invocation is programmed as if the target object was local to the current JVM.

Java RMI Principle




15

The general principle of Java RMI is illustrated in this figure.

A client object in one JVM (left) includes a reference to a server object (remote) in another JVM (right).

This reference is actually a reference to a local stub object (client stub). This stub transforms a method call into a request message (which includes an object reference to identify the object in the server JVM, an method identifier and the parameters of the call). This request message is received by a skeleton object (server stub) which performs the actual method call on the server object.

Java RMI Server side



16

We describe the general functioning the RMI before describing its programming model.

We assume that a Server class has been programmed following the RMI programming model.

On the server side, when the Server class is instantiated, stub and skeleton objects are instantiated. The skeleton object is associated with a local port of the machine for receiving requests.

In order to make the Server object accessible from clients, it must be registered in a naming service called rmiregistry (the rmiregistry runs in another JVM). This registration is possible thanks to the Naming class which provides a bind method (which registers the association between a name ("foo") and the Server object).

This registration in the rmiregistry makes a copy of the stub in the rmiregistry (and registers its association with "foo"). The rmiregistry is ready to deliver copies of the stub to clients.

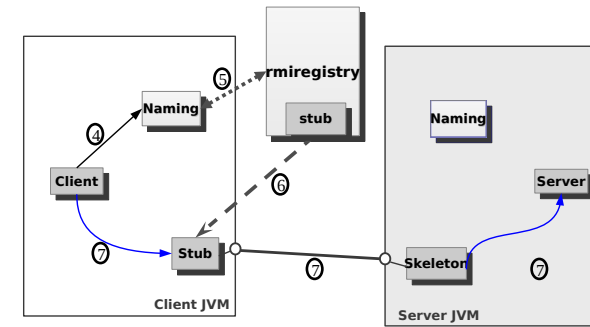
Java RMI Server side

- 0 - At object creation time, a *stub* and a *skeleton* (with a communication port) are created on the server
- 1 - The server registers its instance with a naming service (*rmiregistry*) using the *Naming* class (*bind* method)
- 2 - The naming service (*rmiregistry*) registers the *stub*
- 3 - The naming service is ready to give the *stub* to clients

17

We recall here the main step of creation of a Server object and registration in the *rmiregistry*.

Java RMI Client side



18

On the client side, the client can fetch a reference to the Server object from the *rmiregistry*. This is possible thanks to the *Naming* class which provides a *lookup* method (which queries the object registered with a name ("foo")).

The query on the *rmiregistry* returns a copy of the *stub* (associated with "foo"). This *stub* implements the same interface as the *Server* object. It can be used by the client to invoke a method. The *stub* creates and sends a request message to the *skeleton* which performs the actual call on the *Server* object.

Java RMI Client side

- 4 - The client makes a call to the naming service (*rmiregistry*) using the *Naming* class to obtain a copy of the stub of the server object (*lookup* method)
- 5 - The naming service delivers a copy of the *stub*
- 6 - The *stub* is installed in the client and its Java reference is returned to the client
- 7 - The client performs a remote invocation by calling a method on the *stub*

19

We recall here the main step of querying the *rmiregistry* and performing a method invocation.

Java RMI Utilization

- Coding
 - Writing the server interface
 - Writing the server class which implements the interface
 - Writing the client which invokes the remote server object
- Compiling
 - Compiling Java sources (*javac*)
 - Generation of *stubs* et *skeletons* (*rmic*)
 - (not required anymore, dynamic generation)
- Execution
 - Launching the naming service (*rmiregistry*)
 - Launching the server
 - Launching the client

20

Here are the main steps for using RMI.

Regarding coding :

- you must define the Java interface of the Server. This interface is used both by the Server and the Client.

- the Server class implements the previous interface. The Server is instantiated and the instance is registered in the *rmiregistry*.

- the Client can declare a variable whose type is the previous interface. The Client obtains a copy of the stub from the *rmiregistry*. The stub implements the interface. The Client can call a method on this stub.

Regarding compiling :

- the application is compiled with *javac* as usually

- the stub and skeleton classes can be generated with *rmic* (a stub generator). This is not necessary anymore on recent versions of Java, the stubs being generated dynamically when needed.

Regarding execution :

- you have to launch the *rmiregistry*

- then you can launch the server and then the client

Java RMI Programming

- Programming a remote interface
 - public interface
 - interface: extends java.rmi.Remote
 - methods: throws java.rmi.RemoteException
 - serializable parameters: implements Serializable
 - references parameters: implements Remote
- Programming a remote class
 - implements the previous interface
 - extends java.rmi.server.UnicastRemoteObject
 - same rules for methods

21

Programming RMI applications comes with programming constraints.

For the interface of the Server :

- the interface must be public
- the interface must implement the Remote interface
- all the methods must throw RemoteException
- parameters of remote methods can be of built-in type (int, char), or a Java reference. In this last case, their type must be an interface which is either Serializable or Remote (this is detailed later).

For the Server class :

- it must implement the previous interface
- it must extend the UnicastRemoteObject class
- same rules for methods (as in the previous interface)

Java RMI Example: interface

```
file Hello.java  
  
public interface Hello extends java.rmi.Remote {  
    public void sayHello()  
        throws java.rmi.RemoteException;  
}
```

Description
of the
interface

22

We review a very simple example.

Here is the definition of the interface.

Interface Hello implements Remote and throws RemoteException.

Java RMI Example: server

```
file HelloImpl.java

import java.rmi.*;
import java.rmi.server.UnicastRemoteObject;

public class HelloImpl extends UnicastRemoteObject
    implements Hello {
    String message;

    // Constructor implementation
    public HelloImpl(String msg) throws java.rmi.RemoteException {
        message = msg;
    }
    // Implementation of the remote method
    public void sayHello() throws java.rmi.RemoteException {
        System.out.println(message);
    }
    ...
}
```

Implementation
of the
server class

23

Here is the code of the server class.

Class HelloImpl extends UnicastRemoteObject and implements interface Hello.

Your constructors must throw RemoteException.

The remote method sayHello() throws RemoteException.

Java RMI Example: server

```
file HelloImpl.java

...

public static void main(String args[]) {
    try {
        // Create an instance of the server object
        Hello obj = new HelloImpl("hello");
        // Register the object with the naming service
        Naming.rebind("//my_machine/my_server", obj);
        System.out.println("HelloImpl " + " bound in registry");
    } catch (Exception exc) { ... }
}
}
```

Implementation
of the
server class

NOTICE : in this example, the naming service (rmiregistry)
must have been launched before execution of the server

24

The rest of the code of the server.

The main method creates an instance of the server class (HelloImpl) and registers it in the rmiregistry, thanks to the Naming class.

The URL passed in the rebind() method is `//<machine-name>:<port>/<name>`

- machine-name is the name of the machine which runs the rmiregistry
- port is the port used by the rmiregistry (the default port is 1099)
- name is the name identifying the registered object in the rmiregistry

In its implementation in Java, the rmiregistry has to be colocated (on the same machine) with the JVM which runs the server object. A work around is to implement another rmiregistry (allowing remote registrations).

Notice that after the registration, this is the end of the main method and the JVM would exit. This is not the case, since when we instantiated the server object, a skeleton was instantiated with creation of a communication socket and of a thread waiting for incoming requests. Because of that thread, the JVM does not exit.

Here, we assume that the rmiregistry was launched (rmiregistry is an executable) on the same machine as the server object, with the command :

`rmiregistry <port>` (default is 1099)

Java RMI

running the rmiregistry within the server JVM

```
file HelloImpl.java
public static void main(String args[]) {
    int port;    String URL;

    try {
        Integer l = new Integer(args[0]); port = l.intValue();
    } catch (Exception ex) {
        System.out.println(" Please enter: java HelloImpl <port>"); return;
    }

    try {
        // Launching the naming service - rmiregistry - within the JVM
        Registry registry = LocateRegistry.createRegistry(port);

        // Create an instance of the server object
        Hello obj = new HelloImpl();

        // compute the URL of the server
        URL = "://" + InetAddress.getLocalHost().getHostName() + ":" +
            port + "/my_server";
        Naming.rebind(URL, obj);
    } catch (Exception exc) { ... }
}
```

25

In this other version, we launch a rmiregistry in the same JVM as the one hosting the server object.

The createRegistry method launches a rmiregistry within the local JVM on the specified port.

The interest of doing so is that when you start the application, a rmiregistry is automatically launched and when you kill the JVM, the rmiregistry is also killed. This is very convenient when debugging.

Java RMI Example: client

```
file HelloClient.java
import java.rmi.*;

public class HelloClient {
    public static void main(String args[]) {
        try {
            // get the stub of the server object from the rmiregistry
            Hello obj = (Hello) Naming.lookup("//my_machine/my_server");
            // Invocation of a method on the remote object
            obj.sayHello();
        } catch (Exception exc) { ... }
    }
}
```

Implementation
of the
client class

26

Here is the code on the client side.

It first requests a reference to the target object from the rmiregistry, using the lookup method from the Naming class (the used URL is the same as before).

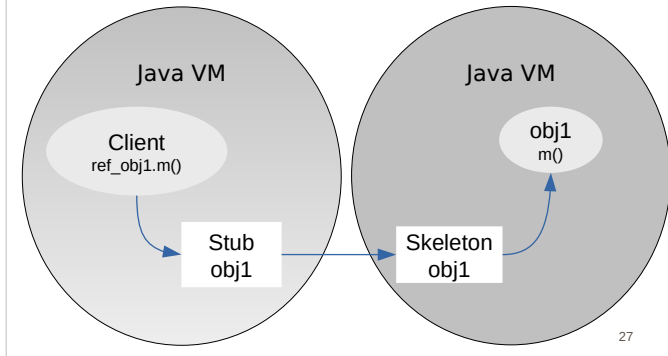
Notice that here the client can be executing on a different machine.

The rmiregistry returns a stub instance. This stub instance implements the same interface as the server object (here Hello). So we can cast the obtained reference with the Hello interface.

Then, invoking a method on the remote object is programmed as if the object was local.

Java RMI


Principle of remote method invocation



To summarize the functioning of Java RMI, a client which obtained (from the rmiregistry) a remote reference (ref_obj1) to a remote object (obj1) has actually a reference to a local stub object (Stub obj1). The client can invoke a method m() on the remote object. It will invoke this method on the stub, which will send the request message. This message is received by the skeleton (Skeleton obj1) which performs the actual invocation on the server object.

Java RMI

Serializable object parameter passing

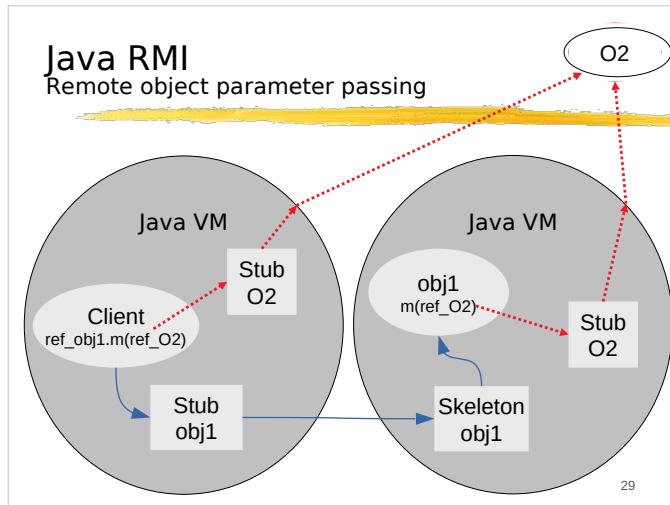


Parameters passed in a remote method can be of built-in types (int char ...). Then the parameters are simply copied (transferred) in the remote server.

If a parameter is a Java reference to an object, the type of the parameter in the method signature must be an interface. Then, there are 2 possibilities:

- Serializable. If the interface is serializable (inherits from Serializable), then the passed object is copied to the server (the object is cloned).
- Remote. If the interface is Remote (inherits from Remote), then the remote reference (i.e. the stub) is passed to the server, meaning that the stub is copied in the server. Therefore, the passed object becomes accessible remotely in the server.
- if the interface is neither Serializable nor Remote, this is an error (it should not compile).

This figure illustrates the Serializable case. The client passes as parameter a reference to object O2 which is local in the client. Then, O2 is copied to the server and the invoked method (m) receives a reference to a clone of object O2 in the server.



This figure illustrates the Remote case. The client passes as parameter a reference to object O2 which is remote (in another JVM). It means that the reference to O2 in the client is a local reference to a sub of O2. Then, the stub of O2 is copied to the invoked server and the invoked method (m) receives a reference to a copy of stub of O2 in the server. Therefore, m() receives a remote reference to O2.

Java RMI Compiling

- Compiling the interface, the server and the client
 - `javac Hello.java HelloImpl.java HelloClient.java`
- Generation of stubs (*not needed anymore*)
 - `rmic HelloImpl`
 - *skeleton* in `HelloImpl_Skel.class`
 - *stub* in `HelloImpl_Stub.class`

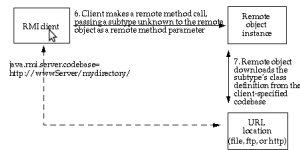
30

To execute the application, you must first compile the interface and the server and client classes.

As previously mentioned, generating stubs and skeletons is not necessary anymore, but you can still do it.

Java RMI Deployment

- Launching the naming service
 - rmiregistry &
- launching the server
 - java HelloImpl
 - java -Djava.rmi.server.codebase=http://my_machine/...
 - URL of a web server from which the client JVM will be able to download missing classes
 - Example: serialization
- Launching the client
 - java HelloClient



Here we explicitly launch the rmiregistry in a shell.
Then, we can launch the server and then the client.

One tricky issue is the availability of classes. Assume the client invokes a method `m(Data d)` on the server, `Data` being an interface which is `Serializable`. Both the client and the server know the interface `Data` (it was necessary to use the method `m` and to compile the code). Then the client may invoke `m` passing an instance of class `ClientData` (which implements `Data`). But the server which receives a copy of the object does not have the `ClientData` class (and different clients may have different implementations of the `Data` interface).

More generally, a JVM may transfer copies of objects (with `Serialization`) to other JVMs. How can the first JVM make these classes available to other JVMs. The solution is to specify, when launching a JVM, a web site from which classes can be downloaded. When classes are missing for using a serialized object, the classes are downloaded and installed dynamically.

```
java -Djava.rmi.server.codebase =URL <a class>
```

When launching a JVM this way, we specify that if serialized instances are given to other JVMs, the missing classes can be found on the web site defined by URL.

Java RMI: conclusion

- Very good example of RPC
 - Easy to use
 - Well integrated within Java
 - Java reference parameter passing: serialization or remote reference
 - Deployment: dynamic loading of serializable classes
 - Designation with URL

Many tutorials about RMI programming on the Web ...
Example : <https://www.javatpoint.com/RMI>

32

To conclude this lecture, Java RMI is a very example of RPC integrated in a Java.

Many tutorials about Java RMI can be found on the net.

Message Oriented Middleware

Daniel Hagimont

<https://www.google.fr/search?q=daniel+hagimont+home+page>

1

This lecture is about messaging services that are provided in the context of Message Oriented Middleware (MOM).

Message based model Introduction

- Client-server model
 - Synchronous calls
 - Appropriate for tightly coupled components
 - Explicit designation of the destination
 - Connection 1-1
- Message model
 - Asynchronous communication
 - Anonymous designation (e.g.: announcement on a newsgroup)
 - Connection 1-N

2

For the moment, we can consider that the message model consists in programming distributed applications with simple message exchanges.

The message model has fundamentally different properties compared to the client-server model.

The client-server model :

- relies on synchronous calls (with a request and a response, the client being suspended waiting for the response)
- is well suited for tightly coupled components, i.e. the caller depends on the service provided by the callee
- there's an explicit designation of the callee by the caller
- it's a one to one connection

In opposition, the message model :

- relies on asynchronous communications (the sender does not wait for a response)
- there can be an anonymous designation (when you send a message to anybody who may be interested like an announcement on a newsgroup)
- it's can be a one to many connection

Message based model Introduction

- Application example
 - Supervision of equipments in a network
 - E.g. average load on a set of servers
- Client-server solution
 - Periodic invocation
- Message based solution
 - Each equipment notifies state changes
 - Administrators subscribe notifications

3

We give here an example of application where the message model is better suited. Let's consider the supervision of equipments in a cluster (e.g. the load of the cluster's machines).

A client-server based solution would require a central server performing periodic invocations of all the servers in the cluster.

A message based solution would see each server notify the central server whenever the load changes.

Message based services ... used everyday

- Electronic forums (News)
 - Pull technologies
 - consumers can subscribe to a forum
 - producers can publish information in a forum
 - consumers can login and read the content anytime
 - Electronic mail
 - Push technologies
 - mailing lists (multicast - publish/subscribe)
 - consumers can subscribe a mailing list
 - producers can send emails to a mailing list
 - Consumers receive emails without having to perform any specific action
- Asynchronous
 - Anonymous
 - 1-N

Motivations : provide communication facilities for developing distributed applications, with such features

4

The message model is already used for many applications in daily use.

For instance, electronic forums (news) are relying on the message model. Producers publish (send) information on a forum. Consumers subscribe to a forum and read (pull) the information published on the forums they subscribed.

Another example is electronic mail with mailing lists. A producer can send email to a mailing list and consumers can subscribe mailing lists and the messages sent (push) to these mailing lists are received by those consumers.

In both examples, communication is asynchronous, anonymous and may involve several receivers.

Message based middleware Principles

- Message Passing (communication with messages)
- Message Queuing (communication with persistent message queues)
- Publish/Subscribe (communication with subscriptions)
- Events (communication with callbacks)

5

Message based middleware were designed to provide developers with a system support for managing messages and programming distributed applications which exchange messages, with the properties presented previously (asynchronous, anonymous, 1-N).

In this context, we distinguish 3 kinds of such messaging service :

- Message passing
- Message queuing
- Publish/subscribe

And one additional service commonly found which is event programming.

Message based middleware Message passing

- Communication with message
 - In a classical environment: sockets
 - In a parallel programming environment: PVM, MPI
 - Other environments: ports (e.g. Mach)


6

Message passing is the simplest service which consists in allowing to send asynchronous messages.

In a classical environment, message passing relies on the socket interface, but it can have other forms, e.g. in an environment devoted to parallel applications (PVM and MPI are parallel environments providing message passing interfaces). Other systems may provides message passing with an interface different from socket (e.g. ports in Mach).

Message based middleware Message Queuing

- Queue of messages
 - persistent messages (reliability)
- Independence between the emitter and the receiver
 - The receiver is not necessarily active
 - => increased asynchronism
 - Several receivers (anonymous)



7

Message Queuing is the first advanced service which may be provided by a MOM.

The basic difference with message passing is that message queuing provides message persistence.


A queue may be allocated and used by clients (producers) or servers (consumers). The queue is managed in the network, meaning that it is not managed in clients neither servers. It is instead managed on machines managed in the middle, i.e. by the message middleware.

Messages are persistent in the sense that we don't require the producer and the consumer to be active at the same time for sending a message (which is the case for message passing). The client may send a message in the queue while the server is inactive (the machine is down). The message will be read by the server at a later time, and may be the client will be inactive at that time.

Another aspect of independence is the fact that a queue may be shared by several producers and consumers. It already provides a sort of anonymous designation.

Message based middleware Publish/Subscribe

- Anonymous designation
 - The receiver subscribes to a topic
 - Subject-based
 - Content-based
 - The producer sends a message to a topic
- Communication 1-N
 - Several receivers may subscribe



8

The second advanced service is the publish/subscribe (pub/sub) service.

A receiver may subscribe to a topic.

There are generally 2 types of pub/sub system:


- subject-based : topics are predefined subjects (i.e. subjects have to be created by an administrator)
- content-based : topics are filters on the content of messages (e.g. I want to receive messages which include)

A producer sends a message to a topic, i.e. either to a given subject or simply with a content. All the receivers who subscribed to the subject, or requested a content which fits with the sent message, will receive a copy of the message.

Here the pub/sub communication service allows message persistence, anonymous designation and multiple receivers.

Message based middleware Events

- Basic concepts: events, reactions (handling associated with event reception)
- Attachement: association between an event type and a reaction



- Exists for all forms of messaging (Message Passing, Message Queuing, Publish/Subscribe)


9

In order not to have to periodically consult message queues (associated with message queuing or pub/sub) message based middleware often introduces support for event programming.

It mainly allows the association between an event (reception of a message) and a reaction (handling program).

Such a facility is available for all forms of communication (message passing, queuing or pub/sub).

Message based middleware Implementations



10

Different implementation strategies may be used.

The simplest one is a centralized server remotely used by all clients. This is appropriate for testing, but not for real use as it represents a single-point-of-failure.

Another organization is an interconnection of distributed servers. The interconnection generally depends on the geographic and administrative distribution of clients. The server may implement routing of messages according to the subscriptions from clients.

The last organization is the software bus where all servers know each others. This is generally a strategy used on local (small scale) networks.

Java Message Service

- JMS: Java API defining a uniform interface for accessing messaging systems
 - IBM (WebSphere MQ), Oracle (WebLogic)
 - Apache ActiveMQ, RabbitMQ
- Message Queue
- Publish/Subscribe
- Event


11

With the popularity of MOMs, and the development of Java, was proposed a common specification of an API for using a MOM from Java. It should be the same API for all messaging systems (from different providers).

This is JMS for Java Message Service. JMS defines a set of Java interfaces which allows a client to access a messaging system. JMS tries to minimize the concepts to learn and manipulate to use a messaging system, while preserving the diversity of all the existing MOMs.

JMS defines interfaces for managing message Queues and Publish/Subscribe.

JMS: an interface (portability, not Interoperability)



Interoperability : AMQP (Advanced Message Queuing Protocol)

12

It is important to note that JMS is an interface. Since it is implemented by many MOM providers, it implies that if you implement your applications with JMS, it will run on many MOMs (from different providers). So JMS addressed the issue of the portability of applications.

However, JMS does not bring interoperability. The messages emitted by provider X may have a different format from those emitted by provider Y.

Portability was brought to MOMs with the standardization of AMQP which defines format of exchanged data at the network level.


JMS interface

- *ConnectionFactory*: factory to create a connection with a JMS server
- *Connection*: an active connection with a JMS server
- *Destination*: a location (source or destination)
- *Session*: a single-thread context for emitting or receiving
- *MessageProducer*: an object for emitting in a session
- *MessageConsumer*: an object for receiving in a session
- Implementations of these interface are specific to providers ...

13

JMS may appear complex, but it is rather systematic, and also it had to satisfy all the providers (if the designers wanted all the providers to implement it).

JMS - Architecture



14

This figure illustrates how these interfaces can be used.

JNDI is the interface of a naming service (such as rmiregistry which is an instance of such a naming service). We assume a JNDI service is available.

A JMS client can obtain from the JNDI service a reference to a *ConnectionFactory*, which allows to create a *Connection* (with the JMS server) and then to create a session in this JMS server.

The JMS client can also obtain from the JNDI service a reference to a *Destination* (an abstract type which can actually refer to a Queue or a Topic).

From a session and a destination, we can create a *MessageProducer* and a *MessageConsumer* allowing to emit and receive messages.

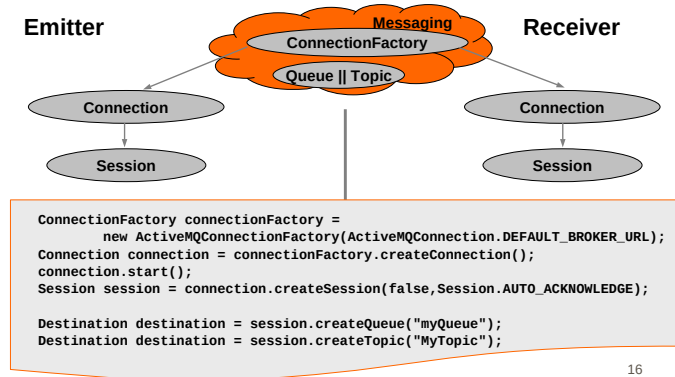
Interfaces PTP et P/S

	Point-To-Point	Publish/Subscribe
<i>ConnectionFactory</i>	<i>QueueConnectionFactory</i>	<i>TopicConnectionFactory</i>
<i>Connection</i>	<i>QueueConnection</i>	<i>TopicConnection</i>
<i>Destination</i>	<i>Queue</i>	<i>Topic</i>
<i>Session</i>	<i>QueueSession</i>	<i>TopicSession</i>
<i>MessageProducer</i>	<i>QueueSender</i>	<i>TopicPublisher</i>
<i>MessageConsumer</i>	<i>QueueReceiver</i>	<i>TopicSubscriber</i>

15

The interfaces described previously are abstract and are specialized according to the use of message queuing (Point-To-Point) or Publish/Subscribe

JMS - initialization



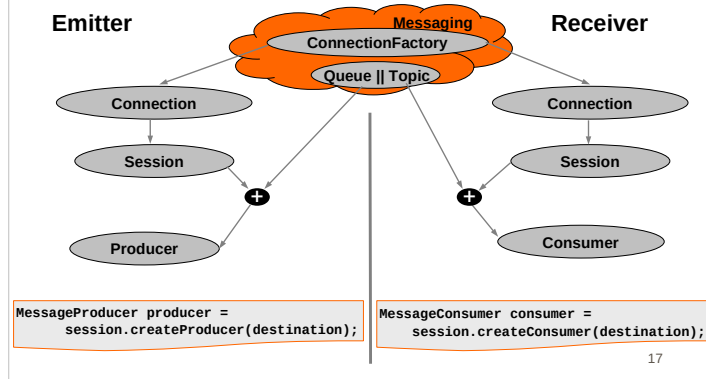
16

Here is the code which is common to the emitter and receiver for initializing the connection with the JMS server and obtaining a destination (one of the 2 lines should be chosen, queue or topic ...).

Notice that with ActiveMQ (this is not JMS standard, but specific to ActiveMQ), createQueue() and createTopic() take a URL as parameter, so the same URL used by 2 clients implies the same destination. These ActiveMQ methods correspond to the query of JNDI.

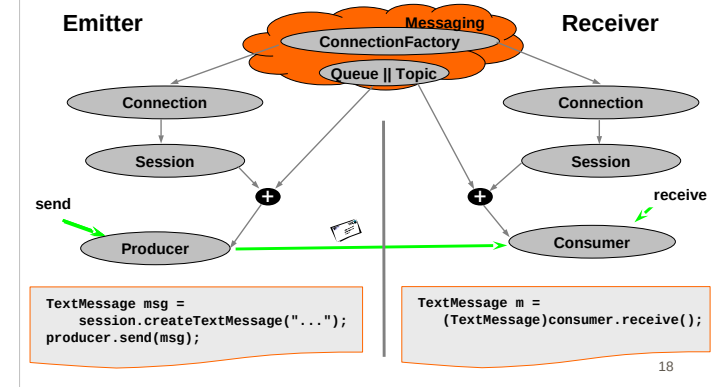
In ActiveMQ, destinations are instantiated at first use.

JMS - producer / consumer



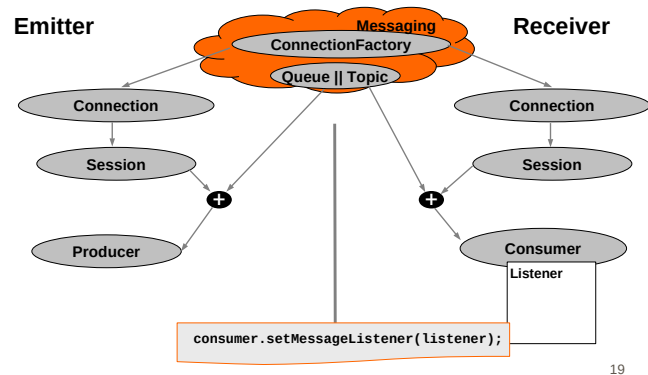
Here, with a session and a destination, we create a producer (left) and a consumer (right).

JMS - communication




On the left, we can send a message (here a `TextMessage`) with a producer. On the right, we can receive a message (here a `TextMessage`) with a consumer.

JMS - Listener



On the consumer side, we can associate a reaction to a message reception event.

JMS - Listener



The registered listener is an instance of a class which implements the `onMessage()` reaction method.

JMS - messages

- TextMessage (a character string)

```
String data;  
TextMessage message = session.createTextMessage();  
message.setText(data);
```

```
String data;  
data = message.getText();
```

- BytesMessage (bytes array)

```
byte[] data;  
BytesMessage message = session.createBytesMessage();  
message.writeBytes(data);
```

```
byte[] data;  
int length;  
length = message.readBytes(data);
```

21

JMS - messages

- MapMessage (sequence of key-value pair)

- A value is a primitive type

```
MapMessage message = session.createMapMessage();  
message.setString("Name", "...");  
message.setDouble("Value", doubleValue);  
message.setLong("Time", longValue);
```

```
String name = message.getString("Name");  
double value = message.getDouble("Value");  
long time = message.getLong("Time");
```

22

In JMS, messages are types. We can allocate :

- TextMessage (like String)
- BytesMessage (like byte[]).

JMS - messages

- StreamMessage (sequence of values)
 - A value is a primitive type
 - Reading should respect the sequence order to writing

```
StreamMessage message = session.createStreamMessage();  
message.writeString("...");  
message.writeDouble(doubleValue);  
message.writeLong(longValue);
```

```
String name = message.readString();  
double value = message.readDouble();  
long time = message.readLong();
```

23

JMS - messages

- ObjectMessage (serialized objects)

```
ObjectMessage message = session.createObjectMessage();  
message.setObject(obj);
```

```
obj = message.getObject();
```

24

Conclusions

- Communication with messages
 - Simple programming model
 - Many extensions, variants ...
 - Message software bus, actors models, multi-agent systems
 - ...
 - Widely used for interconnecting tools, existing, developed independently
- However... it is only apparently simple
 - Propagation and report of errors
 - Development tools
- Tutorials
 - <https://www.jmdoudoux.fr/java/dej/chap-jms.htm>

25

Even if the message model may seem to be very simple and primitive, many extensions and variants exist.

MOMs are widely used for interconnecting tools, integrating tools that were developed independently.

Notice that simplicity is only apparent, as asynchronism makes it difficult to debug or to have deterministic behaviors.

Web Services

Daniel Hagimont

<https://www.google.fr/search?q=daniel+hagimont+home+page>

1

This lecture is about web services.

Motivations

- Motivations
 - Coarse-grained application integration
 - Unit of integration: the "service" (interface + contract)
- Constraints
 - Applications developed independently, without anticipation of any integration
 - Heterogeneous applications (models, platforms, languages)
- Consequences
 - Elementary common basis
 - For communication protocols (messages)
 - For the description of services (interface)
 - Base choice: HTTP and XML/JSON

2

The example of RPC tool we have seen, Java RMI, is restricted to interactions within Java applications, allowing remote invocations of Java objects.

With Web services, the motivation is to provide a RPC facility for the interaction (and integration) of coarse-grained applications (that we call services). A service is supposed to be much bigger than a simple Java object.

Web services aim at allowing the interaction between application developed independently, with different environments (models, platforms, languages).

Web services rely on elementary existing protocols and formats: mainly HTTP, XML and JSON.

Basic form of WS : XML-RPC (1998)

Description in XML of a remote procedure call
Parameter types are specified in an XML schema

```
<methodCall>
  <methodName>meteo.temperature</methodName>
  <params>
    <param>
      <value><int>31130</int></value>
    </param>
  </params>
</methodCall>
```

Description in XML of parameter returns

```
<methodResponse>
  <params>
    <param>
      <value><int>25</int></value>
    </param>
  </params>
</methodResponse>
```

Interest : independence with respect to
platforms and communication protocols

<https://en.wikipedia.org/wiki/XML-RPC>

3


XML-RPC was a precursor of what are web services now.

XML-RPC was a RPC protocol relying of XML for the representation of requests and HTTP for the transport of requests.

The idea was to be independent from execution platforms or languages and to rely on widely recognized and adopted formats.

XML-RPC was a precursor and evolved into SOAP, the protocol used in web services.

SOAP WS : architecture (2000++)



4

This figure illustrates the architecture of web services (WS).

A service provider may implement a service in any language and/or platform, as soon as a runtime for WS exists in his environment.

The runtime is a composed of

- stub and skeleton generators
- a WSDL generator
- a web server (WS runtime) for making services available on the internet

Then, on the server side, the service implementation is linked with the web server, in order to be able to receive requests through the HTTP communication protocol. A skeleton is generated and is a web application in the web server. A WSDL description (Web Service Description Language) of the service is generated and published, i.e. made available to potential clients.

The WS architecture specifies that a service registry (a naming service) should be used for the publication and discovery of WSDL descriptions. However, UDDI was not actually used. Generally the WSDL of a WS can be published on a Web server as any document.

On the client side, the WSDL description can be copied and used to generate a stub in the environment of the client. Notice that the environment of the client is not mandatorily the same as the one of the server. Then the client can implement an application which is able to invoke the WS by calling the stub.

The stub communicates with the skeleton with the SOAP/HTTP protocol which is a standard.

HTTP and SOAP are standards from the W3C.

SOAP describes the syntax of request and response messages which are transported with HTTP.

Elements of SOAP WS

- Description of a service
 - WSDL : Web Services Description Language
 - Standard notation for the description of a service interface
- Access to a service
 - SOAP : Simple Object Access Protocol (over HTTP)
 - Internet protocol allowing communication between Web Services
- Registry of services
 - UDDI : Universal Description, Discovery and Integration
 - Protocol for registration and discovery of services

5

Therefore the main elements of WS are :

- the description of the service in WSDL. Generally, from an implementation of a service (e.g. a procedure), tools are provided to generate the WSDL description of the service, which is published for clients. The clients can use this WSDL description to generate stubs so that calls to the service can be programmed easily.
- access protocols which are SOAP (for the content of messages) and HTTP (for the transport). All the WS runtimes (in any environment) comply with these standards.
- registries of service (UDDI) which are not really used.

Tools

- From a program, we can generate a WS skeleton
 - Example: from a Java program, we generate
 - A servlet which receives SOAP/HTTP requests and reproduces the invocation on an instance of the class
 - A WSDL file which describes the WS interface
- The generated WSDL file can be given to clients
- From WSDL file, we can generate a WS stub
 - Example: from a WSDL file, we generate Java classes which can be used to invoke the remote service
- Programming is simplified
- Such tools are available in different language environments

6

To illustrate this, we give an example of use in the Java environment.

In the Java environment, a WS tool is used to generate from a program (with an exported interface) a skeleton as a servlet. A servlet is a Java program which runs in a web server. This servlet/skeleton receives SOAP/HTTP requests and reproduces the invocation on an instance of the class. The WSDL specification of the WS is also generated.

The WSDL file is published on the web and imported by the client.

From the WSDL specification, the client can generate stubs which make it easier to program WS invocations.

In the following slides, we give an example with Apache Axis.

Example: programming a Web Services

- Eclipse JEE
- Apache Axis
- Creation of a Web Service
 - From a Java class
 - In the Tomcat runtime
 - Generation of the WSDL file
- Creation of a client application
 - Generation of stubs from a WSDL file
 - Programming of the client

7

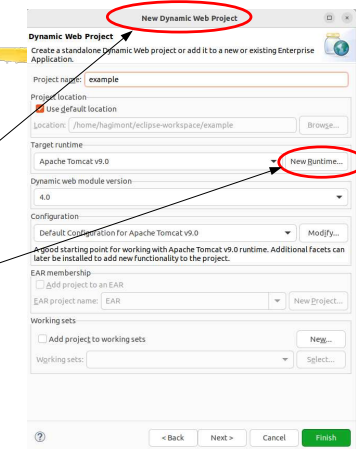
We use Eclipse JEE and Apache Axis which is available in Eclipse JEE.

Apache Axis is used to generate from a Java class a servlet which is installed in the Tomcat engine (the web server). It also generates the WSDL description which describes the interface of the WS.

On the client side, the WSDL description is used to generate stubs which are used to invoke the WS in a client program.

Create a Dynamic Web Project


- Eclipse JEE
- Open JEE perspective
- Create a Dynamic Web Project
- Add your Tomcat runtime



8

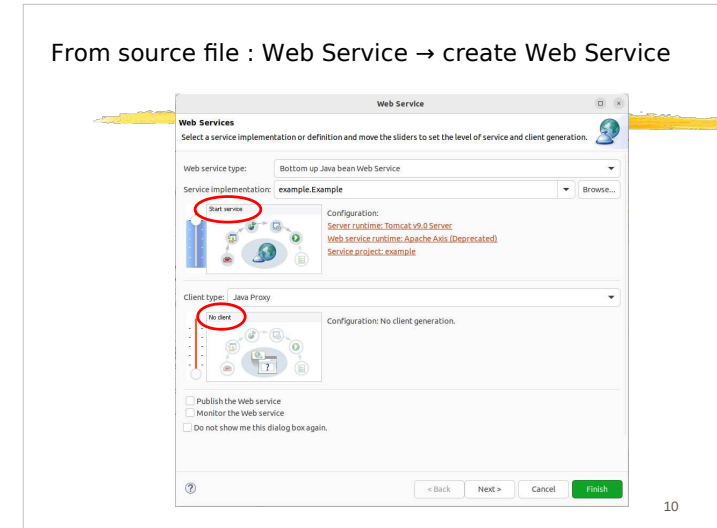
In Eclipse, we create a dynamic web project (a project allowing the develop servlets) and add the Tomcat runtime.

Create a Class



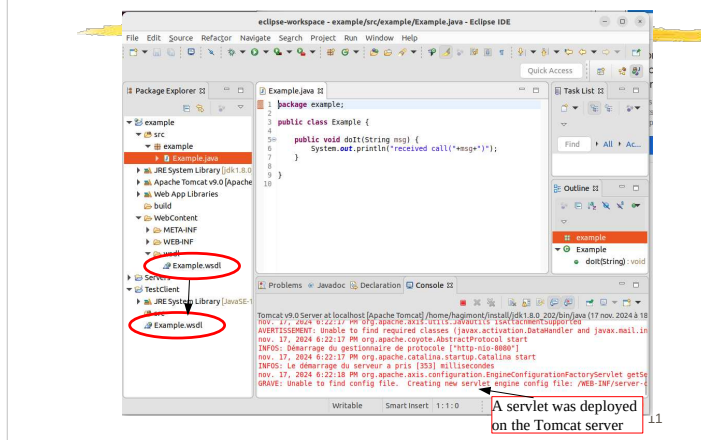
In this project, we create a class.
Notice that a Tomcat server is running in Eclipse.

From source file : Web Service → create Web Service



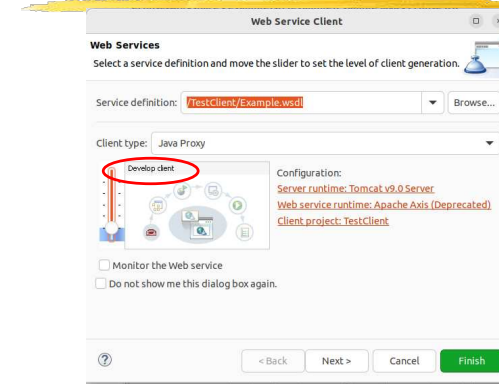
From the source file of the class, we can generate (right click) a WS from this file.

Copy the generated WSDL file in a new Java project



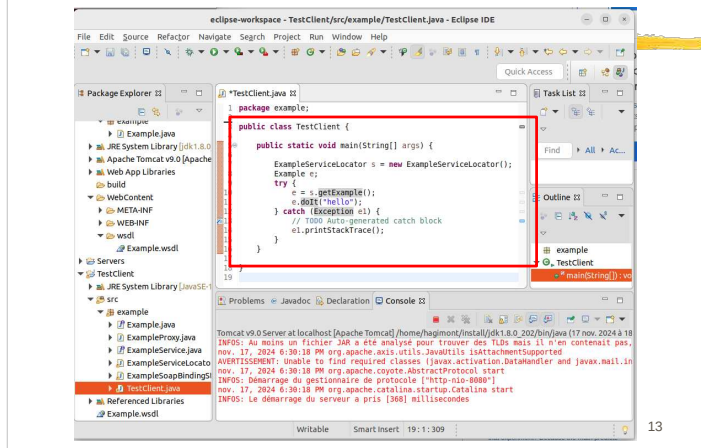
Then, we create a new Java project and copy the WSDL description in the new project.

From the WSDL file
Web Service → Generate Client (Develop Client)



In the new project, from the WSDL file, we generate (right click) the stubs (develop Client).

Program a client

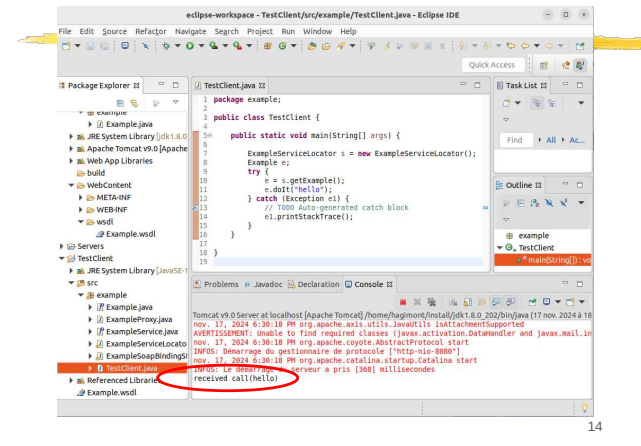


13

In the new project, we can program an application which makes an invocation of the WS.

The procedure to follow to invoke the WS depends on the tool used (here Apache Axis).

Run



14

We can then run the client program which invokes the WS.

Generated WSDL

```
<wsdl:definitions targetNamespace="http://DefaultNamespace"
xmlns:apachesoap="http://xml.apache.org/xml-soap" xmlns:impl="http://DefaultNamespace"
xmlns:intf="http://DefaultNamespace" xmlns:wsdl="http://schemas.xmlsoap.org/wsdl/"
xmlns:wsoap="http://schemas.xmlsoap.org/wsdl/soap/"
xmlns:xsd="http://www.w3.org/2001/XMLSchema">
<!--WSDL created by Apache Axis version: 1.4
Built on Apr 22, 2006 (06:55:48 PDT)-->
<wsdl:types>
<schema elementFormDefault="qualified" targetNamespace="http://DefaultNamespace"
xmlns="http://www.w3.org/2001/XMLSchema">
<element name="dolt">
<complexType>
<sequence>
<element name="msg" type="xsd:string"/>
</sequence>
</complexType>
</element>
<element name="doltResponse">
<complexType/>
</element>
</schema>
</wsdl:types>
```

15

We can have a look at the WSDL description.

We can see that the WSDL syntax is not very simple. Therefore such WSDL descriptions are not written by the user, but generally generated by the tool on the server side and imported by the client.

Generated WSDL

```
<wsdl:message name="doltResponse">
<wsdl:part element="impl:doltResponse" name="parameters">
</wsdl:part>
</wsdl:message>
<wsdl:message name="doltRequest">
<wsdl:part element="impl:dolt" name="parameters">
</wsdl:part>
</wsdl:message>
<wsdl:portType name="MyService">
<wsdl:operation name="dolt">
<wsdl:input message="impl:doltRequest" name="doltRequest">
</wsdl:input>
<wsdl:output message="impl:doltResponse" name="doltResponse">
</wsdl:output>
</wsdl:operation>
</wsdl:portType>
```

16

Very verbose !

Generated WSDL

```
<wsdl:binding name="MyServiceSoapBinding" type="impl:MyService">
  <wsdlsoap:binding style="document" transport="http://schemas.xmlsoap.org/soap/http"/>
  <wsdl:operation name="doIt">
    <wsdl:operation soapAction=""/>
    <wsdl:input name="doItRequest">
      <wsdlsoap:body use="literal"/>
    </wsdl:input>
    <wsdl:output name="doItResponse">
      <wsdlsoap:body use="literal"/>
    </wsdl:output>
  </wsdl:operation>
</wsdl:binding>
<wsdl:service name="MyServiceService">
  <wsdl:port binding="impl:MyServiceSoapBinding" name="MyService">
    <wsdlsoap:address location="http://localhost:8080/HW/services/MyService"/>
  </wsdl:port>
</wsdl:service>
</wsdl:definitions>
```

17

Very very verbose !

SOAP request(with TCP/IP Monitor)

```
<?xml version="1.0" encoding="UTF-8"?>
<soapenv:Envelope xmlns:soapenv="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">

  <soapenv:Body>
    <doIt xmlns="http://DefaultNamespace">
      <msg>hello</msg>
    </doIt>
  </soapenv:Body>

</soapenv:Envelope>
```

18

We can have a look at the SOAP request. This is simply a standardized format for exchanged messages.

SOAP response

```
<?xml version="1.0" encoding="utf-8"?>
<soapenv:Envelope xmlns:soapenv="http://schemas.xmlsoap.org/soap/envelope/"
  xmlns:xsd="http://www.w3.org/2001/XMLSchema"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">

  <soapenv:Body>
    <doltResponse xmlns="http://DefaultNamespace"/>
  </soapenv:Body>

</soapenv:Envelope>
```

19

Here is the SOAP response.

REST Web Services (2000+ +)

- A simplified version, not a standard, rather a style (of use of simple features of the web)
- Increasingly popular
- Use of HTTP methods
 - GET : get data from the server
 - POST : create data in the server
 - PUT : update data in the server
 - DELETE : delete data in the server
- Invoked service encoded in the URL
 - http://<machine>/module/service
 - For instance
 - HTTP request to <machine>
 - GET /module/service
 - service only returns data from the server

20

SOAP/WSDL based WS were very popular some years ago. They are becoming obsolete.

An evolution of WS is REST WS. This is a simplified version which is very popular now. Notice that REST WS is not a standard, but rather a recommendation or a style of implementation, based on simple features of the web.

It relies on HTTP requests (GET, POST, PUT, DELETE), but mainly GET and POST are used. GET is used when you want to read (only) data from the WS while POST is used when you want to modify something in the WS.

The service that you call is encoded in the URL.

For instance a GET HTTP request on URL http://<machine>/module/service

This service is supposed to return data from the server, without updating anything in the server (its a style, it is not enforced).

REST Web Services

- **Parameter passing**
 - HTTP parameters
 - Format : field1=value1&field2=value2
 - GET : parameters in the URL
 - `http://nom_du_serveur/cgi-bin/script.cgi?champ1=valeur1&...`
 - POST : parameters are included in the body of the HTTP request
 - XML or JSON (or any other)
 - Included in the body of the HTTP request
- **Response**
 - Included in the body of the HTTP response: XML or JSON (or any other)
- **Description of a REST WS (available on the net)**
 - A simple document describing the methods and parameters
 - No WSDL

21

Parameter passing can be based on HTTP parameters. With HTTP parameters, parameters are encoded as a String `field1=value1&field2=value2` ...

This parameter string is passed in the URL if you use the GET HTTP method, and it is passed in the body of the request if you use the POST HTTP method.

You can also pass parameters in a document (XML or JSON or any format) in the body of the request.

A service can return a document (XML or JSON) in the body of the response (which corresponds to return parameters).

The description of a REST WS is simply a document describing the services that you may call and the passed parameters (names, formats).

Example of existing REST WS

Currency API Request
The base URL for our currency API is
`https://www.amdoren.com/api/currency.php`

Request Parameters

Parameter	Description
api_key	Your assigned API key. This parameter is required.
from	The currency you would like to convert from. This parameter is required.
to	The currency you would like to convert to. This parameter is required.
amount	The amount to convert from. This parameter is optional. Default is a value of 1.

Example:
To get the latest exchange rate in EUR for 1 USD:
`https://www.amdoren.com/api/currency.php?api_key=1BZdLwH2yCvXjgTZ6x&from=USD&to=EUR`

Currency API Response

Element	Description
error	Error code. Value greater than zero indicates an error. See list below.
error_message	Short description of the error. See list below.
amount	The exchange rate or amount converted.

Example:
JSON data returned from our currency API request:
{ "error" : 0, "error_message" : "", "amount" : 0.90368 }

22

Here is an example of description of a REST WS. This is for a currency converter.

It says that you have one service available :

<https://www.amdoren.com/api/currency.php>

It lists the parameters that may be passed in the HTTP GET request. A example is given.

It then describes the response which is a JSON. A example is given.

Development of REST WS

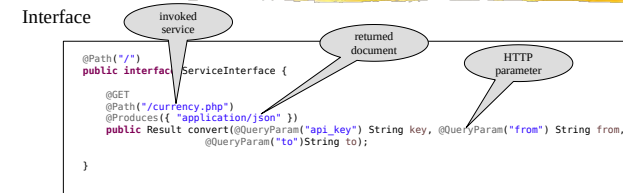
- You can develop your application by hand in any programming language
 - Verbose and error prone
 - As for RPC, code can be generated
- Many development environments
 - e.g. Resteasy and Jersey
 - Resteasy in the following of the talk (client and server sides)
 - A view on Spring (mainly server side)

23

Many development environments can be used to REST WS development.

In the following, we overview the used of resteasy (on the server side and the client side) and we have a look at Spring. Both will be used in the labs.

Existing REST WS (client with resteasy)



Java bean (generated from JSON)

```
public class Result {

    String error;
    String error_message;
    String amount;

    // getters/setters

}
```

24

With the currency converter, as said in the documentation, the conversion method takes 3 HTTP parameters (api_key, from, to, the last is optional) and it returns a JSON.

The 3 HTTP parameters are associated with Java parameters (with @QueryParam) and a Java bean is created for the JSON.

Existing REST WS (client with resteasy)

Client

```
public class Client {  
    public static void main(String args[]) {  
        final String path = "https://www.andoren.com/api";  
        ResteasyClient client = new ResteasyClientBuilder().build();  
        ResteasyWebTarget target = client.target(UriBuilder.fromPath(path));  
        ServiceInterface proxy = target.proxy(ServiceInterface.class);  
        Result r = proxy.convert("9xwjRjxTtnzuGKH7LcWC5Vengr52F3", "EUR", "AMD");  
        System.out.println("convert: "+r.getError()+"-"+r.getError_message()+  
            +"/"+r.getAmount());  
    }  
}
```

easy
invocation

25

And here is an example of client which invokes the service.

Implementing a service with resteasy

WS class

```
@Path("/")  
public class Facade {  
    static Hashtable<String, Person> ht = new Hashtable<String, Person>();  
    @POST  
    @Path("/addperson")  
    @Consumes({ "application/json" })  
    public String addPerson(Person p) {  
        ht.put(p.getId(), p);  
        return "person added";  
    }  
    @GET  
    @Path("/getperson")  
    @Produces({ "application/json" })  
    public Person getPerson(@QueryParam("id") String id) {  
        return ht.get(id);  
    }  
    @GET  
    @Path("/listpersons")  
    @Produces({ "application/json" })  
    public Collection<Person> listPersons() {  
        return ht.values();  
    }  
}
```

Receives a JSON
Deserialized into a Java object
Returns a String

Returns an object
Serialized into a JSON
Receives an id HTTP parameter

Person is a simple POJO

26

As for SOAP/WS, many tools were implemented to help developers.

Here, we present Resteasy (Jersey is also a very popular one you may look at).

On the server side, you can use annotations in a Java program to say :

- each method is associated with a path in the URL used to access the WS
- `@Path` : specifies the element of the path associated with the class or the method. Here method `addPerson()` is associated with path `/addperson`
- `@POST` or `@GET` : specifies which HTTP method is used. Notice that GET returns an object (data) while POST returns an HTTP code (and a message).
- `@Consumes` : specifies that we receive a JSON object which is deserialized into a Java object.
- `@Produces` : specifies that we return a Java object which is serialized into a JSON object.
- `@QueryParam` : the `getPerson()` method has an "id" parameter. The `QueryParam` annotation associates this parameter with an "id" HTTP parameter.

Implementing a service with resteasy

- Add the RestEasy jars in Tomcat (lib)
- In eclipse (not easy with vscode)
 - Create a Dynamic Web Project
 - Add RestEasy jars in the buildpath
 - Create a package
 - Implement the WS classes (Facade + Person)
 - Add a class RestApp

```
public class RestApp extends Application {  
    private Set<Object> singletons = new HashSet<Object>();  
    public RestApp() {  
        singletons.add(new Facade());  
    }  
    public Set<Object> getSingletons() {  
        return singletons;  
    }  
}
```

27

To run this example :

- add the Resteasy jars in Tomcat and Eclipse
- create a dynamic web project (a servlet project)
- add the RestEasy jars in the buildpath of the project
- create a package and the previously developed classes
- add the RestApp class

Implementing a service with resteasy

- Add a web.xml descriptor in the WebContent/WEB-INF folder

```
<?xml version="1.0" encoding="UTF-8"?>  
<web-app xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"  
    xmlns="http://xmlns.jcp.org/xml/ns/javaee"  
    xsi:schemaLocation="http://xmlns.jcp.org/xml/ns/javaee  
        http://xmlns.jcp.org/xml/ns/javaee/web-app_3_1.xsd" version="3.1">  
    <display-name>essai-server</display-name>  
    <servlet>  
        <servlet-name>resteasy-servlet</servlet-name>  
        <servlet-class>  
            org.jboss.resteasy.plugins.server.servlet.HttpServletDispatcher  
        </servlet-class>  
    </servlet>  
    <init-param>  
        <param-name>javax.ws.rs.Application</param-name>  
        <param-value>pack.RestApp</param-value>  
    </init-param>  
    <servlet-mapping>  
        <servlet-name>resteasy-servlet</servlet-name>  
        <url-pattern>/*</url-pattern>  
    </servlet-mapping>  
</web-app>
```

- Export the war in Tomcat

28

Add the web.xml descriptor in the WebContent/WEB-INF folder and export a war (in the webapps folder of Tomcat)

Publish the WS

- Just write a documentation which says that
 - The WS is available at <http://<machine-name>:8080/<project-name>/>
 - Method addperson with POST receives a person JSON :

```
{
  "id": "00000",
  "firstname": "Alain",
  "lastname": "Tchana",
  "phone": "0102030405",
  "email": "alain.tchana@enseeiht.fr"
}
```

- Method getperson with GET receives an HTTP parameter id and returns a person JSON
 - Method listperson returns a JSON including a set of persons
- A caller may use any tool (not only RestEasy)

29

Publication of a REST WS is simply a document describing the interface.

Implementing the client with resteasy

- From the previous documentation, a client can write the interface

```
@Path("/")
public interface FacadeInterface {

    @POST
    @Path("/addperson")
    @Consumes({ "application/json" })
    public String addPerson(Person p);

    @GET
    @Path("/getperson")
    @Produces({ "application/json" })
    public Person getPerson(@QueryParam("id") String id);

    @GET
    @Path("/listpersons")
    @Produces({ "application/json" })
    public Collection<Person> listPersons();
}
```

30

On the client side, from the documentation, a user can write a Java interface with Resteasy annotations. Of course, it's very similar to what we wrote on the server side, but we could do it for a WS we don't know (we only have the documentation).

Implementing the client with resteasy

- And write a class which invokes the WS

```
public class TestClient {  
    public static void main(String[] args) {  
        final String path = "http://localhost:8080/rest-server";  
        ResteasyClient client = new ResteasyClientBuilder().build();  
        ResteasyWebTarget target = client.target(UriBuilder.fromPath(path));  
        FacadeInterface proxy = target.proxy(FacadeInterface.class);  
  
        String resp;  
        resp = proxy.addPerson(new Person("007", "James Bond"));  
        System.out.println(resp);  
        resp = proxy.addPerson(new Person("006", "Dan Hagi"));  
        System.out.println(resp);  
  
        System.out.print("List Person: ");  
        Collection<Person> l = proxy.listPersons();  
        for (Person p : l) System.out.print("["+p.getId()+"/"+p.getName()+"]");  
        System.out.println();  
  
        Person p = proxy.getPerson("006");  
        System.out.println("get Person: "+p.getId()+"/"+p.getName());  
    }  
}
```

31

The previous annotated interface (FacadeInterface) makes it easy to invoke the service. We can build a proxy object of type FacadeInterface.

This proxy allows programming service invocations simply as method calls.

Implementing the client with resteasy

- In eclipse or vscode
 - Create a Java Project
 - Add RestEasy jars in the project
 - Implement the Java bean that correspond to the JSON
 - Automatic generation with <https://www.site24x7.com/tools/json-to-java.html>
 - Implement the interface and the client class (FacadeInterface + TestClient)
 - Run

32

This is the procedure to run the client.

A view on Spring

- A development environment for server side
 - Spring-boot: facilitate the configuration
 - Relies on Maven (dependencies)
 - Can produce
 - A standalone application (including Tomcat)
 - A war to be deployed in Tomcat
 - Also provides client side support (within a Spring server)
 - In VSCode
 - Extension: Spring initializr java support

33

Spring is a development environment for developing REST WS.

Spring-boot is an extension which simplifies the configuration. It relies on Maven.

Thanks to Spring-boot, you can produce:

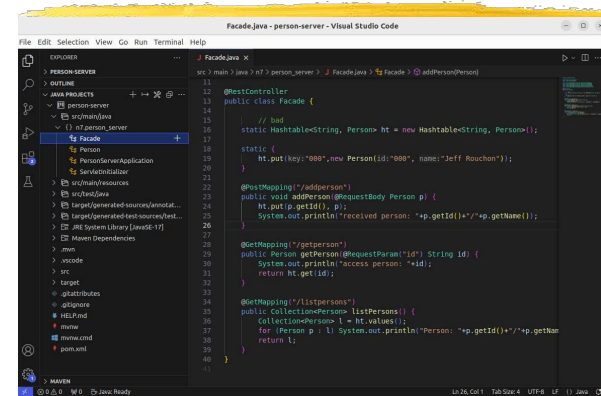
- an application which embeds Tomcat. When you launch it, you start a Tomcat web server which included your REST WS application.

- a war archive which includes your REST WS application. This war can be deployed in a running Tomcat server.

Spring also provides support for invoking an external REST WS from a Spring WS.

In VSCode, you can use the "Spring initializr" extension which automates the creation of a Spring-boot project.

Spring initializr : Create a Maven Project Implement an annotated class



Assuming you have installed in VSCode the Spring initializr java support extension.

Here, we create the same REST WS as with RestEasy (person management).

You can create a Spring initializr project.

Initially, there are 2 classes in the project (PersonServerApplication and ServletInitializer). You don't have to modify them.

We just implement a class (Facade) which implements the methods of the REST WS.

You can annotate these methods with @GetMapping and @PostMapping.

Method parameters can be annotated with @RequestParam (for HTTP parameters) or @RequestBody (for a JSON object).

By default, the returned object is serialized into a JSON object.

35

If you run: `./mvnw spring-boot:run`

It will start a Tomcat server which

You can test your WS with a web browser.

Notice that the URL (<http://localhost:8080>)

The screenshot shows a web browser window with a REST client interface. The address bar displays 'localhost:3000/person'. The main content area shows a JSON response for a person with ID 1. The response is highlighted in blue and includes the following details:

```
{
  "id": 1,
  "name": "John",
  "age": 30,
  "address": {
    "street": "123 Main St",
    "city": "New York",
    "state": "NY",
    "zip": "10001"
  },
  "email": "john.doe@example.com",
  "phone": "123-456-7890"
}
```

The interface also includes a sidebar with tabs for 'JSON', 'Headers', 'Cookies', and 'Status'. The 'JSON' tab is currently selected. The status bar at the bottom indicates the response status is 200.

36

In your VScode project, you can run: `./mvnw package`

It will produce a war in the target folder (person-server-0.0.1-SNAPSHOT.war).

Here, I copy this war into my Tomcat (webapps folder).

I rename it as `person-server.war` as the name of the war is used in the URL used

The URL (<http://localhost:8080/person-server/listpersons>) used in the browser

includes the name of the war.

Invoking a rest API from Spring

```
11 @RestController
12 public class ClientController {
13     final static String path = "http://localhost:8080/person-server/";
14
15     @GetMapping("/call")
16     public Collection<Person> call() {
17
18         try {
19             RestClient client = RestClient.create();
20             client.post().uri(path + "addperson").contentType(MediaType.APPLICATION_JSON).body(new Person(){{"id": "111", "name": "Don Mac"}}).retrieve().toBodylessEntity();
21             Collection<Person> list = client.get().uri(path + "listpersons").retrieve().body(new ParameterizedTypeReference<>());
22             return list;
23         } catch (Exception ex) {
24             ex.printStackTrace();
25             return null;
26         }
27     }
28 }
```

37

Spring also provides support for invoking external REST WS from within a Spring WS.

The class to use for that is `RestClient`.

Here, I programmed a simple Spring REST WS with a method "call". In that method, I invoke methods "addperson" and "listpersons" from the previous service.

Conclusion

- Web Services: a RPC over HTTP, exchanging XML or JSON
- Interesting for heterogeneity as there are tools in all environments
- Recently
 - SOAP WS less used
 - REST + XML/JSON more popular
 - Micro-services: used for structuring backend applications

38

To conclude, Web services aim at implementing a RPC service on top of HTTP and relying on standard formats (XML, JSON).

One of the main interest is the independence between the server (the service provider) and the client (the service consumer). They can be from different organizations and use different tools, OS, or languages.

The recent evolution is an obsolescence of SOAP and an increased popularity of REST and JSON.

Recently, the micro-service architecture was proposed. It consists in architecting large applications (especially backend applications) in terms of a set of interconnected REST WS (components). The advantage is independence between components.